Source Coding - Final Project

The LZ77 and LZSS Algorithms for Data Compression

Tommaso Martini (108 15 80)

July 25, 2014

1   Introduction

This project has the aim of developing a dictionary-based compression coding technique and comparing the gained performances with those of softwares present in a typical PC. In particular, we are required to develop our own version of the LZ77 and LZSS algorithms.

In the following we will give a brief overview of such algorithms and we will explain how our versions have been implemented, focusing on the implementation choices. Eventually, we will provide a comparison between the two algorithms and their commercial versions, when they work to compress several kinds of file.

1.1   Dictionary-based coding

A *dictionary-based coding* is a coding technique which processes a file as a sequence of symbols to build a dictionary, that is a sequence of pairs (key, value), where the value can be either a symbol or a sequence itself. The dictionary, properly formatted, is the message to send to the receiver. The latter has to build back the original sequence exploiting the information contained in the dictionary entries. Actually, for the dictionaries we will be using, we do not care very much about the keys, since the entries are stacked with the same order the receiver will use at its side, therefore we will not talk about a key for a dictionary entry, rather than about a position.

There are two kinds of dictionary-based coding technique, depending on the dictionary form. The first is the *implicit dictionary coding*, which is the case of the LZ77 algorithm, in which the entries are generated and read sequentially, without the need of specifying a key for them. On the other hand, the *explicit dictionary coding* creates some entries which are referred to with their keys and, in the decoding process, are used also nonsequentially. This is the case of the LZ78 algorithm, published by the same authors of the LZ77 one year later.

2   The LZ77 Algorithm

In this section we provide a short overview of the dictionary-based compression algorithm presented by Abrham Lempel and Jacob Ziv in 1977 [4]. The aim of this algorithm is trying to reduce the redundancy due to similar sequence repeated along the message. What it does is basically checking if a certain piece of the sequence has already been found in the past and, if it is the case, we code the whole piece with a reference to the previous alike sequence.

2.1   The algorithm

The algorithm works with two adjacent windows shifting on the right during the coding of the message. The first window is the *searching*

window and has length L_s; the second one is the *coding window* and has length L_c. We want to find a match between any prefix of the *coding window* and a sequence contained in the overall window given by the juxtaposition of the *searching window* and the *coding window*, i.e. it is sufficient that only the first symbol of the match belongs to the *searching window*, while the matched pattern can stretch in through the *coding window* itself. The window lengths L_s and L_c are two parameters of the algorithm and their choice strongly influences the compression performances, we will see how in the following part of this report.

Once we have found a match between a prefix of the *coding window* and another piece of message, we code it as a triplet (*offset*, *length*, *symbol*) where:

- **offset** denotes the number of back hops we have to do, starting from the first symbol of the *coding sequence*, to find the beginning of the matched string; it is clear that, for how it is defined, the *offset* cannot exceed the *searching window* dimension;

- **length** is the length of the matched string or, equivalently, of the prefix of the *coding window* which has been matched. This value cannot exceed L_c, since this is the maximum dimension of the pattern we are looking for;

- **symbol** is the first symbol after the matched prefix of the *coding window*. Its encoding assures the functioning of the algorithm also in case of no matches.

After the generation and the storage of a triplet inside the dictionary, we shift the windows such that the first symbol of the *coding window* corresponds to the first non encoded symbol.

2.2   LZ77 Implementation

In the following paragraphs we will present our own implementation in Matlab of the algorithm,

spanning the different versions we produced. As a matter of fact, several editions of the same program have been written, adopting different approaches in some parts of the code. Actually the versions differ above all for the pattern matching algorithm that has been adopted, while the rest of the code remains basically the same.

2.2.1   Basic Coding Algorithm

First of all we choose a file and write it as a sequence of bytes (`uint8` in Matlab). We assume this is the original message, whose alphabet is made of all the possible combinations of 8 bits and has size $M = 2^8 = 256$. The first symbol of the sequence is always encoded as a single symbol, because placing the *coding window* in the first positions makes no sense, since we would have no *searching window* to exploit. Then, the *coding window* starts from the second position of the sequence and, until there are at least L_s symbols before it, the *searching window* is shrunk to fit into the available positions. The same thing happens to the *coding window* when we are reaching the end of the file: if less than L_c symbols are left, we use a shorter *coding window*. Note that, in order to always have a symbol to insert in the *symbol* field of the last triplet, near the end of the file the *coding window* is shrunk such that the last symbol of the stream is left out.

Once the dictionary has been created, we want to write it using the least amount of memory we can. Since the values of L_s and L_c are fixed and they determine the maximum values for the fields *offset* and *length* and we know that a symbol is represented with a single byte, we can evaluate the maximum number of bits to use for each component of the triplet as:

$$l_{offset} = \lceil log_2 L_s \rceil \text{ bits} \tag{1}$$

$$l_{length} = \lceil log_2 L_c \rceil \text{ bits} \tag{2}$$

We use exactly $l_{offset} + l_{length} + 8$ bits for each triplet. After creating a sequence of bits by lining

up all the rows of the dictionary, we can split it in blocks of 8 bits and code it as a sequence of bytes. To let the receiver know how many bits represent each element of the triplet, we attach in front of the stream three bytes: the first indicating the number of bits related to the *offset*, the second related to the *lenght* and the third to the *symbol*. One may argue that, with a single byte, we can referring only to values smaller than $2^8 = 256$; this is true, but, for our purposes, a byte is more than sufficient: suppose we need 200 bits to represent the *offset* in a triplet; this means that our *searching window* is almost 2^{200} bits long, that is around 2×10^{50} Gigabytes, of course absolutely useless for ordinary files. At the receiver side the user inspects the first three bytes to retrieve the number of bits dedicated to each component of the triplet, then, after looking at the message as a stream of bits, he can group together the triplets of the dictionary by exploiting the information just got. This explains also why we cannot use less than 3 bytes to express the dimension of one triplet: since the *byte* is the fundamental unity of memory in every computer, we are sure that every user is able to get the initial information and thus properly decode the dictionary.

2.2.2   Version 1: Basic

As previously mentioned, the many versions of the algorithms differ by the pattern matching technique to find the *coding window* inside the *searching window*. The very basic version of the program executes the pattern matching by brute force: we compare the first character of the *coding window* with each element of the *searching window*. When a match occurs, we compare the second element of the *coding window* with the following symbol of the *searching window* and so on until we find two different symbols. Then, if the length of the match is larger than the longest match found so far, we replace it and save also the starting point of the match in the *searching window*. If, in the end, the longest match has

zero length, this means that no match has been found and we have to encode one only symbol. This basic algorithm is very wasteful, because we have to try every possible position.

2.2.3   Version 2: KMP

A first improvement has been brought by using the KMP (Knuth-Morris-Pratt) pattern matching algorithm [1]. This algorithm is optimum for the research of a pattern inside a string: it takes a time $\mathcal{O}(n+k)$, where n is the length of the pattern and k is the length of the string. We thought about using the KMP algorithm where the pattern **c** is the *coding window* and the overall string $[\mathbf{s}, \mathbf{c}]$ is the *searching window* concatenated with the *coding window*.

The KMP algorithm starts with the brute force approach, comparing the first character of the pattern with each character of the string, but it allows to improve time performances in pattern matching because, thanks to previous failed matches, it realizes when it is possible to jump ahead and discard some trials which will surely fail. We used the "jump mechanism" of KMP to realize a second version of the program, but the time performances keep remaining very poor, due to the many nested cycles we have to use, which are not very fast to be executed in Matlab. Moreover, each different pattern needs a *matching function* to be computed, which increases the number of necessary computations.

2.2.4   Version 3: optimistic `strfind()`

The third version uses the Matlab function `strfind()` for pattern matching, which is optimized to work in Matlab. We simply invoke this method using the *searching window* as string and the *coding window* as pattern to be found. If no matches occur we discard the last symbol of the *coding window* and repeat the procedure until either a match is found or we reduce the length of the pattern to zero, that means that we have to

encode the symbol individually.

The term *optimistic* refers to the order in which we take different patterns: we are confident to find the longest match, the best one, quite soon and therefore we start by looking for the entire pattern. Of course, if there is little redundancy and there is often no match found, we have to repeat the cycle L_c times for each different *coding window*.

2.2.5   Version 4: pessimistic `strfind()`

From an empirical point of view we found that we often have not many long matches and hence the *optimistic* approach is not very suitable. The *pessimistic* approach consists in starting to look for a match with the shortest possible prefix of the *coding window* and executing the pattern matching through `strfind()` until we find no more matches. In this case we consider the previous match, which is the longest one, to build the current triplet. If there are no matches we know it from the first iteration and this allows us to save a lot of time. Quite all the results presented in the following have been obtained with this version of the program.

2.3   Further improvements

A further improvement we thought about, but we did not implemented, is using a sort of bisectional approach to find the longest match: we first try to search the first character of the *coding window*; if we do not find it, this means that there is no other matches, otherwise we look for the entire look-ahead buffer. Again, if a match with the whole *coding window* is found, we stop searching, otherwise we try with the first two symbols of the buffer and so on. The idea is keeping searching a longer buffer at one step, starting with length 1, and a shorter buffer at the successive step, starting with the entire *searching window*. This method is worse than the previous

ones if we already know whether the file is higly or poorly redundant, but, in case of a completely unknown file, it allows to save about half of the computational time. Since in our experience processing time was not a strong constraint and we perfectly know what kind of files we are dealing with, we did not realize this implementation, choosing, time by time, the most suitable version between the optimistic and pessimistic approach.

Another possible enhancement we could bring to the algorithm is checking what is the most convenient match to encode in a triplet inside the same look-ahead buffer. It may happen that, even if we find a prefix of the *coding window* corresponding to a past symbol sequence, by encoding it we lose the capability of detecting another longer match in the following. A trivial, but expensive, solution to this issue is searching the best match among all the possible subsequences of the *coding window*.

3   The LZSS Algorithm

The LZ77 algorithm has a lack of performances due to the constraint on the triplet size: whenever we encode a single symbol, we waste $l_{offset} + l_{length}$ bits because the *offset* and the *length* of the match are both zero, in fact the triplet to send is $[0, 0, \alpha]$. In 1982 Storer and Szymanski [3] proposed a different version of the LZ77 algorithm which improves the compression performances through a better management of the size of the dictionary entries.

We can distinguish two kinds of entries: those which contain the information related to an only symbol not find in the *searching window* and those which refer to a previous pattern. We can use a flag bit to indicate in which case we are. If the bit is 0 the entry is made only by the code of a symbol. If the bit is 1 the entry is made of a pair (*offset*, *length*). This convention allows us to save a lot of space in both cases.

The encoding of a single symbol implies the

employment of 9 bits: 8 for the symbol itself and one for the 0 flag. The encoding of a pair, instead, takes $l_{offset} + l_{length} + 1$ bits. It is clear that it could be more convenient to encode some symbols independently even if a match has been found, if the match length is shorter than:

$$k = \left\lceil \frac{l_{offset} + l_{length} + 1}{9} \right\rceil \qquad (3)$$

3.1 LZSS Implementation

The main difference with the LZ77 implementation lies on the dictionary construction, all the rest of the code is quite the same used in the implementation of the LZ77 exposed in section 2.2.5, with a pessimistic usage of the function strfind() for pattern matching.

Once again, to save memory, we first look at the dictionary as a stream of bits lining up its rows, then we write it as a sequence of bytes and, eventually, we attach in front of it three bytes expressing the number of bits employed to encode the fields *offset*, *length* and *symbol*.

4 Performances of the Algorithms

We want now to focus our attention on the performances reached by our implementations of the LZ77 and LZSS algorithms in relation both with the type of file we are trying to compress and the parameters setting, that is the lengths of the *searching window* and the *coding window*.

4.1 Testing files

All the experiments we performed have been executed on seven specific types of files, part of them taken from the *Canterbury Corpus*, a collection of files with several formats to be used as sample files for compression testing. The used files are here reported:

- **rep.txt**: simple text file containing a sequence of character A repeated 10000 times (10000 bytes);

- **per**: file containing 10 repetitions of a random sequence of 1000 bytes (10000 bytes);

- **ran**: file generated by concatening 10000 random bytes (10000 bytes);

- **shak.txt**: simple text file reporting a piece of a Shakespeare poem (13741 bytes);

- **web.html**: html format page (24603 bytes);

- **code.c**: piece of c code (11150 bytes);

- **sum**: SPARC executable file (38240 bytes).

We chose these files to have a wide span among different kinds of files and show, in this way, the several behaviours the algorithm can assume for different scenarios. In particular, we use the file rep.txt to obtain the maximum compression capability of the algorithm, the file per was included to show the behaviours in presence of periodicity and the file ran represents the worst scenario for compression purposes, since it has no redundancy to exploit and a high entropy. The last four files, on the other hand, do not represent borderline or special cases, but more common examples of files one can deal with during his ordinary life.

4.2 Redundant file

The first file, rep.txt, is the most redundant file we can have: 10000 repetitions of the same character, A. The nature of the file allows a very strong compression, since we could just say what is the symbol and how many times it is repeated. The following table shows the performances of the algorithms executed with a fixed *coding window* of length 2000 and a *searching window* varying between 1000, 5000 and 10000, which is the length of the message itself:

rep.txt		
L_s	LZ77	LZSS
1000	0.25%	0.19%
5000	0.27%	0.20%
10000	0.28%	0.21%

As we expected the compression is very high: we can represent 10000 bytes with less than 30 bytes. It may look strange that performances drop with the increasing of the *searching window*; this is due to the number of bits we need to encode a triplet, which raises with the length of the window (in LZ77 we use 29 bits when $L_s = 1000$, 32 bits when $L_s = 5000$ and 35 bits when $L_s = 1000$). LZSS works to solve this issue, as a matter of fact it behaves slightly better than LZ77 and manages to encode the dictionary with less bits. In this case we can see that the length of the *coding window* has also a great relevance on performances; the followig table summarizes the results obtained when we use the same three searching lengths of the previous example, but a coding length $L_c = 10000$:

rep.txt		
L_s	LZ77	LZSS
1000	0.11%	0.08%
5000	0.22%	0.08%
10000	0.12%	0.09%

In this case we just need about ten bytes to represent the whole message and the dictionary is composed by only two rows. It is interesting to report here the 3D plot of the LZ77 performances, in percentage of compression, in function of L_s and L_c:


Figure 1: LZ77 compression ratio in function of L_s and L_c.

The above plot well shows that for this file the choice of the *coding window* is more important than the choice of the *searching window*, which, we will see, is not common for a generic file.

4.3 Periodic file

A big problem with implicit dictionary based compression algorithms is that periodicity of the message could not be properly exploited. This unfavorable event occurs if the windows are shorter than the period of the message, such that there is no way the algorithm can become aware of the existance of periodicity in the message; the LZ78 algorithm was born to solve this issue. We can bring a clear example compressing the file `per` keeping fixed the *coding window* to 2000, but with two different *searching window* lengths: 500 (half a period) and 5000 (five times a period). The compression results are reported in the following table:

per		
L_s	LZ77	LZSS
500	189.98%	112.53%
5000	23.55%	11.44%

It is clear that, in the first case, the algorithm is not exploiting the periodicity, in fact it is ex-

panding the data and increasing the file dimension, because, being the period built randomly, it finds no redundancy to exploit. In the second scenario, however, periodicity is exploited and the file is strongly compressed.

4.4   Randomly generated file

When a file is made by a random sequence of bytes its entropy is very high and the algorithm finds no redundancy to exploit for compression. In this case performances are very low, in fact the output file has a greater dimension than the input file and we can do very little also by varying L_s and L_c. The following table reports the compression ratios with $L_c = 2000$ and $L_s = 1000,\ 5000,\ 10000$:

ran		
L_s	LZ77	LZSS
1000	184.66%	112.53%
5000	197.83%	112.55%
10000	202.12%	112.56%

We can notice two main facts: the first is that LZSS compresses in general better than LZ77; the second is that LZSS performances decrease much more slowly than LZ77; if we inspect the dictionaries created by the two algorithms, we can try to spot the reason of such difference: the LZ77 dictionary contains a lot of references to an only previous symbol and we cannot avoid using a whole triplet to encode them. On the other hand the LZSS algorithm is able to realize that it is not a good deal to waste a whole triplet to express an only symbol and it encodes it as a single byte. In this way we ideally would have a dictionary with the same dimensions of the message, but, unlucklily, the coding algorithm also includes the flag bits. Thus, it is practically quite impossible to compress a random sequence: if we try to compress it, we are going to hopelessly expand it.

4.5   Common files

The file shak.txt is an example of human-written text file. As we can see from the results of the experiment, LZ77 and LZSS are not the best choice for compressing this kind of file. The reason is that the spoken language has not a strong redundancy and patterns of characters are often short and repeated only a few times. The table below shows the performances with $L_c = 2000$:

shak.txt		
L_s	LZ77	LZSS
1000	89.48%	77.91%
5000	82.26%	72.52%
10000	81.38%	71.93%

Code files as web.html and code.c presents a greater redundancy because of their scientific format; strict rules of the programming languages make them very repetitive and thus the LZ77 and the LZSS manage to perform a better compression of such files. The following tables show the results for the two files where a 2000 symbols long *coding window* has been used:

code.c		
L_s	LZ77	LZSS
1000	73.53%	59.72%
5000	62.25%	52.98%
10000	58.68%	50.14%

web.html		
L_s	LZ77	LZSS
1000	59.82%	49.27%
5000	51.65%	43.02%
10000	51.45%	42.93%

The executable file sum can also be quite well compressed:

sum		
L_s	LZ77	LZSS
1000	68.91%	60.58%
5000	61.85%	56.70%
10000	55.83%	51.61%

To conclude this section it is interesting to plot a surface graph of the performances in function of both L_s and L_c. We compressed files shak.txt and code.c changing both the *searching window* and the *coding window* from $L_s = L_c = 1000$ to the dimension of the files themselves, with step of 1000. Note that, if the windows are larger than the file sizes, it makes no difference how large the buffers are, because the windows are shrunk to adapt to the actual length of the message.

Figure 4: LZ77 compression ratio for code.c.

Figure 2: LZ77 compression ratio for shak.txt.

Figure 5: LZSS compression ratio for code.c.

Figure 3: LZSS compression ratio for shak.txt.

Two main things can be noticed from the above graphs. First of all we have a confirmation that the LZSS works in general better than the LZ77, as a matter of fact plots related to the LZSS lie below the corresponding LZ77 plots. The second, and most relevant information, links the window sizes with the compression performances: it is intuitive that large *searching windows* are more suitable, because, if we consider a long sequence of symbols, finding some repeated patterns is more likely; on the other hand, and this could be less intuitive, increasing the *coding window* is not very advisable: it is clear that a too short

8

coding window is not a good deal, because we risk to encode a unique, long pattern with more pairs (*offset, length*) than necessary. But why should enlarging it be even worse? The reason is that, in common files as the ones we used for the experience, repeated patterns are never very long, therefore having a huge *coding window* does not represent a profit, in fact we need many bits to represent the *length* field. To conclude, even though performances depend on the kind of file we are trying to compress, we should prefer a quite short *coding window* and a large *searching window*, being careful that this choice may affect time performances.

5 Comparison with existing softwares

In this section we report some results obtained by the comparison between our implementations and other softwares usually installed in common PCs. We want to mention right now that our implementations have no chances to behave as well as any of the professional softwares employed by default by a PC. The reasons are mainly two: the first one is that programs that compress files or folders into archives .ZIP use an advanced algorithm, where the simple LZ77 is enhanced by further coding, in particular by a Huffman code; this combined coding is called LZ77 Deflate, is used by most of the .ZIP compressors and will be shortly described in the following. The second reason is about time: since we execute our program in Matlab, performances are lowered by the infrastructure of the environment and, moreover, `for` cycles are not very convenient in this particular language.

5.1 LZ77 Deflate

The *deflate* algorithm differs from the simple LZ77 for the final compression of the dictionary through a Huffman code. The description of the basic version of the algorithm can be found in the RFC1951 [2], here we provide only a brief overview of its functioning.

The original file is first split into blocks of data whose maximum length must be 64 Kbytes. The LZ77 dictionary is generated through the known procedure, with $L_s = 32000$ and $L_c = 258$ bytes. Once the dictionary is generated, it is compressed using a Huffman code; even if we have three types of data to express (*symbol, offset* and *length*), we use only two Huffman trees. The first tree is made of 285 codewords: the first 256 $(0, \ldots, 255)$ stand for literals, that is the symbols (we look at the message as a sequence of bytes), value 256 indicates the end of the block of data and values $257, \ldots, 285$ are used to indicate the *length* field, which can assume values in $0, \ldots, 258$. Note that some of the *length* codewords are used to represent more than one actual length, for example value 277 refers to the set of lengths $67, \ldots, 82$. To properly encode a length it is, thus, necessary to attach after the codeword itself a sequence of bits, at most five, which specifies the *length* value among the set described by the codeword. Each value $257, \ldots, 285$ has a fixed number of following bits (from no one to 5 bits), hence, since Huffman is a prefix code, the receiver can uniquely decode a codeword just by its prefix and then extract the additional information from the following bits without misunderstanding the code. The second Huffman tree is used for the distances between the matched pattern from the beginnning of the *coding window*. Of course, if a *length* is decoded, the receiver knows that the following bits refer to the related *offset*.

Each block of data is preceded by three bits: the first one says whether the current block is the last one; the second and the third specify the type of compression that has been used: not compressed, compressed with fixed and defined Huffman trees or compressed with dynamically generated Huffman trees. If the Huffman code

is generated ad hoc, we must also transmit the tree to the receiver to allow a correct interpretation. In general the same set of probabilities can give birth to different Huffman trees, but, specifying some construction rules, we can guarantee the unique correspondence between alphabet and tree. Moreover, a tree generated in such a way can be encoded using only the sequence of the codeword lengths. The Huffman trees are, in turn, compressed combining Huffman and run-length codes. The description of this code is included in the message itself. For a dynamically compressed block of data, the structure of the message is the following:

- 3 bits to specify whether the block is the last one and how it is compressed;

- description of the code used to compress the Huffman tree for *symbols* and *lengths* and the Huffman tree for the *offsets*;

- compressed Huffman trees for *symbols*, *lengths* and *offsets*;

- compressed representation of the LZ77 dictionary;

- byte 256 to indicate the end of the block.

Figure 6: Highly redundant files.

Figure 7: Common files.

5.2 Comparison with `zip` and `gzip`

Combining several steps of compressions it is possible to reach very good performances, also in cases where simple LZ77 or LZSS are not very effective. In the following bar graphs we compare, for each of the seven test files, the compression performances of our implementation of the LZ77 and LZSS algorithms and the Unix versions of the programs zip and gzip:

## 6 Conclusions

As we expected, for common files and also for a random sequence of bytes (`ran`), the commercial software behaves better than our own implementation, thanks to the different steps of compression that have been applied. Another aspect we could have guessed is that the pure LZ77 algorithm is the worst one. Anyhow, one could won-

der why in the first case, `rep.txt`, our versions are the best ones. This is due to the additive information brought by the LZ77 Deflate algorithm: the Huffman coding is very useful in a generic case, but the overall information needed to decode it could actually be over-necessary if the file is very redundant. On the other hand, we can see that, for a high entropy file as `ran` the commercial software is very robust, even though no ones among the algorithms can go below the 100% compression threshold, which is, actually, impossible from information theory. In this case, however, additional Huffman and run-length coding manage to approach the maximum possible compression. In all the remaining cases, `zip` and `gzip` act better than our programs, but we can note that the general trend is almost the same for all the four programs.

References

[1] J.H. MORRIS D. KNUTH and V. PRATT. Fast pattern matching in strings. *SIAM Journal on Computing*, 1977.

[2] P. DEUTSCH. Rfc1951, deflate compressed data format specification version 1.3, 1996.

[3] J.A. STORER and T.G. SZYMANSKI. Data compression via textual substitution. *J.ACM*, 1982.

[4] J. ZIV and A. LEMPEL. A universal algorithm for sequential data compression. *IEEE Transactions on Information Theory*, 1977.